

## Exercise: Pods

Pods are the smallest, most basic deployable objects in Kubernetes. A Pod represents a single instance of a running process in your cluster. Pods contain one or more containers, such as Docker containers. Although you want to deploy pods directly (static pods), knowledge for defining pods manifest files will be used for defining more complex Kubernetes resources like Controllers.

### Practice1: Simple pods operations

**Note:** Try not to do a copy/paste on commands requests unless you are instructed to do so. Copy/paste will not help you to learn Kubernetes!

1. Login to Azure and connect to your AKS cluster.

We created an Azure Kubernetes Service cluster by navigating to 'Kubernetes Services' and selecting 'Create'. The process for creating the cluster is similar to that for a virtual machine. After creating the cluster, we used the following commands to connect to it from Azure CLI.

```
PS /home/vladimir> az account set --subscription c031718c-8f31-4c8d-ae56-634904ec50f9
PS /home/vladimir> az aks get-credentials --resource-group kuber --name Test
Merged "Test" as current context in /home/vladimir/.kube/config
PS /home/vladimir>
```

2. Check how many pods run under the default namespace. Run **kubectl get pods**.

```
PS /home/vladimir> kubectl get pods
No resources found in default namespace.
```

3. You should not see any pod under the default namespace. Now check all namespaces. Run **kubectl get pods --all-namespaces**.

```
PS /home/vladimir> kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	azure-ip-masq-agent-lvdvx	1/1	Running	0	13m
kube-system	azure-ip-masq-agent-vzkwj	1/1	Running	0	13m
kube-system	cloud-node-manager-85pxp	1/1	Running	0	13m
kube-system	cloud-node-manager-ffklf	1/1	Running	0	13m
kube-system	coredns-59b6bf8b4f-4r1ln	1/1	Running	0	14m
kube-system	coredns-59b6bf8b4f-bvncj	1/1	Running	0	12m
kube-system	coredns-autoscaler-5655d66f64-z7ts2	1/1	Running	0	14m
kube-system	csi-azuredisk-node-6sssk	3/3	Running	0	13m
kube-system	csi-azuredisk-node-c9f2d	3/3	Running	0	13m
kube-system	csi-azurefile-node-gb2cf	3/3	Running	0	13m
kube-system	csi-azurefile-node-z8ssh	3/3	Running	0	13m
kube-system	konnektivity-agent-7f75c5f96f-6xxf4	1/1	Running	0	14m
kube-system	konnektivity-agent-7f75c5f96f-p9ktc	1/1	Running	0	12m
kube-system	kube-proxy-gwd78	1/1	Running	0	13m
kube-system	kube-proxy-wpvnp	1/1	Running	0	13m
kube-system	metrics-server-8655f897d8-lf67v	2/2	Running	0	12m
kube-system	metrics-server-8655f897d8-rtnlz	2/2	Running	0	12m

4. How many pods do you see? Who deployed these pods? Why are they deployed?

There are 17 pods that were deployed automatically by the Azure Kubernetes Service (AKS) when we created our cluster. They are part of the Kubernetes infrastructure and perform various tasks that are essential for the proper functioning of our cluster. Here is a brief description of each of the pods:

- **azure-ip-masq-agent:** This pod is responsible for configuring IP masquerading rules to enable network connectivity between our Kubernetes pods and the external network.
- **cloud-node-manager:** This pod manages the lifecycle of the worker nodes in our cluster and ensures that the correct number of nodes is available at all times.
- **coredns:** This is a DNS server that provides name resolution for our Kubernetes services and pods.
- **coredns-autoscaler:** This pod scales the number of coredns replicas based on the number of nodes in our cluster.
- **csi-azuredisk-node:** These pods provide support for the Azure Disk storage class in our cluster.
- **csi-azurefile-node:** These pods provide support for the Azure File storage class in our cluster.
- **konnnectivity-agent:** This pod provides secure connectivity between our cluster and other Kubernetes clusters or services running on different networks.
- **kube-proxy:** This is a network proxy that runs on each worker node and enables communication between the pods and services in our cluster.
- **metrics-server:** This pod collects resource utilization data (CPU, memory, etc.) for our cluster and provides it to other Kubernetes components and tools.

5. Now deploy your first pod using the imperative approach. Run **kubectl run nginx --image=nginx**.

```
PS /home/vladimir> kubectl run nginx --image=nginx
pod/nginx created
```

6. Validate if the pods has been created. What is the status of your pod?

```
PS /home/vladimir> kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   0           24s
```

7. Check the logs coming out of your pod. Run **kubectl logs nginx**.

```
PS /home/vladimir> kubectl logs nginx
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2023/04/03 21:46:31 [notice] 1#1: using the "epoll" event method
2023/04/03 21:46:31 [notice] 1#1: nginx/1.23.4
2023/04/03 21:46:31 [notice] 1#1: built by gcc 10.2.1 20210110 (Debian 10.2.1-6)
2023/04/03 21:46:31 [notice] 1#1: OS: Linux 5.4.0-1104-azure
2023/04/03 21:46:31 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2023/04/03 21:46:31 [notice] 1#1: start worker processes
2023/04/03 21:46:31 [notice] 1#1: start worker process 29
2023/04/03 21:46:31 [notice] 1#1: start worker process 30
2023/04/03 21:46:31 [notice] 1#1: start worker process 31
2023/04/03 21:46:31 [notice] 1#1: start worker process 32
```

8. Run following command to check current resource consumption of your pod: **kubectl top pod nginx**.

```
PS /home/vladimir> kubectl top pod nginx
NAME      CPU(cores)   MEMORY(bytes)
nginx     0m           5Mi
```

9. Check on which Node your pods has been scheduled. Run **kubectl get pods -o wide**.

```
PS /home/vladimir> kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP          NODE                                     NOMINATED NODE   READINESS GATES
nginx     1/1     Running   0           6m5s  10.244.0.8  aks-agentpool-24052243-vmss000001    <none>           <none>
```

10. Try to find the same information but this time running **kubectl describe pod nginx**.

```
PS /home/vladimir> kubectl describe pod nginx
Name:      nginx
Namespace: default
Priority:   0
Service Account: default
Node:      aks-agentpool-24052243-vmss000001/10.224.0.5
Start Time: Mon, 03 Apr 2023 21:46:28 +0000
Labels:    run=nginx
Annotations: <none>
Status:    Running
IP:        10.244.0.8
```

11. Delete your pod using **kubectl delete pod nginx**.

```
PS /home/vladimir> kubectl delete pod nginx
pod "nginx" deleted
PS /home/vladimir>
```

12. Let's find the image used on one of the coredns pods under the kube-system namespace.

13. Once again list all pods under all namespaces.

```
PS /home/vladimir> kubectl get pods --all-namespaces
NAMESPACE      NAME                                     READY   STATUS    RESTARTS   AGE
kube-system    azure-ip-masq-agent-lvdvx              1/1     Running   0           46m
kube-system    azure-ip-masq-agent-vzkwj              1/1     Running   0           46m
kube-system    cloud-node-manager-85pxp               1/1     Running   0           46m
kube-system    cloud-node-manager-ffklf               1/1     Running   0           46m
kube-system    coredns-59b6bf8b4f-4r1ln               1/1     Running   0           46m
kube-system    coredns-59b6bf8b4f-bvncj               1/1     Running   0           45m
kube-system    coredns-autoscaler-5655d66f64-z7ts2    1/1     Running   0           46m
kube-system    csi-azuredisk-node-6sssk               3/3     Running   0           46m
kube-system    csi-azuredisk-node-c9f2d               3/3     Running   0           46m
kube-system    csi-azurefile-node-gb2cf               3/3     Running   0           46m
kube-system    csi-azurefile-node-z8ssh               3/3     Running   0           46m
kube-system    konnectivity-agent-7f75c5f96f-6xxf4    1/1     Running   0           46m
kube-system    konnectivity-agent-7f75c5f96f-p9ktc    1/1     Running   0           45m
kube-system    kube-proxy-gwd78                       1/1     Running   0           46m
kube-system    kube-proxy-wpvnp                       1/1     Running   0           46m
kube-system    metrics-server-8655f897d8-lf67v        2/2     Running   0           44m
kube-system    metrics-server-8655f897d8-rtnlz        2/2     Running   0           44m
```

14. Note one of the coredns pods. Now run `kubectl describe pod <coredns-name> -n kube-system`. Replace the <coredns-name> place holder with noted name.

```
PS /home/vladimir> kubectl describe pod coredns-59b6bf8b4f-bvncj -n kube-system
```

15. Inspect the output and locate the image information.

```
PS /home/vladimir> kubectl describe pod coredns-59b6bf8b4f-bvncj -n kube-system
Name:                coredns-59b6bf8b4f-bvncj
Namespace:           kube-system
Priority:              2000001000
Priority Class Name:   system-node-critical
Service Account:      coredns
Node:                 aks-agentpool-24052243-vmss000000/10.224.0.4
Start Time:           Mon, 03 Apr 2023 21:12:40 +0000
Labels:               k8s-app=kube-dns
                     kubernetes.io/cluster-service=true
                     pod-template-hash=59b6bf8b4f
                     version=v20
Annotations:          prometheus.io/port: 9153
Status:               Running
IP:                   10.244.1.3
IPs:                  IP: 10.244.1.3
Controlled By:        ReplicaSet/coredns-59b6bf8b4f
Containers:
  coredns:
    Container ID:   containerd://50a15d64ca6b5b4add8f808b2120ea4c601d673ee8c93566586f22aa60de7bcb
    Image:           mcr.microsoft.com/oss/kubernetes/coredns:v1.9.3
    Image ID:        sha256:c38f956b642366c8eeb0babfda6b0bb2aa92f27a968589804cadb445f6df72d6
```

16. Now let us check the logs of the metrics-server pod. Run the same command as in step 7 but don't forget to add the namespace in which this pod is created.

**"kubectl logs metrics-server-8655f897d8-rtnlz -n kube-system -c metrics-server"**

**-c flag** is used to specify the container name

**-n flag** is used to specify the namespace of the pod

This will retrieve the logs only for the **metrics-server** container running in the **metrics-server-8655f897d8-rtnlz** pod in the **kube-system** namespace.

```
PS /home/vladimir> kubectl logs metrics-server-8655f897d8-rtnlz -n kube-system -c metrics-server
I0403 21:13:05.603746 1 serving.go:342] Generated self-signed cert (/tmp/apiserver.crt, /tmp/apiserver.key)
I0403 21:13:11.404567 1 requestheader_controller.go:169] Starting RequestHeaderAuthRequestController
I0403 21:13:11.404590 1 shared_informer.go:240] Waiting for caches to sync for RequestHeaderAuthRequestController
I0403 21:13:11.404598 1 secure_serving.go:266] Serving securely on [::]:4443
I0403 21:13:11.404616 1 configmap_cafile_content.go:201] "Starting controller" name="client-ca:kube-system::extension-apiserver-authentication"
I0403 21:13:11.404637 1 dynamic_serving_content.go:131] "Starting controller" name="serving-cert::/tmp/apiserver.crt::/tmp/apiserver.key"
I0403 21:13:11.404643 1 shared_informer.go:240] Waiting for caches to sync for client-ca:kube-system::extension-apiserver-authentication
I0403 21:13:11.404658 1 configmap_cafile_content.go:201] "Starting controller" name="client-ca:kube-system::extension-apiserver-authentication"
I0403 21:13:11.404662 1 shared_informer.go:240] Waiting for caches to sync for client-ca:kube-system::extension-apiserver-authentication
I0403 21:13:11.404992 1 tlsconfig.go:240] "Starting DynamicServingCertificateController"
W0403 21:13:11.405106 1 shared_informer.go:372] The sharedIndexInformer has started, run more than once is not allowed
I0403 21:13:11.603781 1 shared_informer.go:247] Caches are synced for RequestHeaderAuthRequestController
I0403 21:13:11.603902 1 shared_informer.go:247] Caches are synced for client-ca:kube-system::extension-apiserver-authentication
I0403 21:13:11.705205 1 shared_informer.go:247] Caches are synced for client-ca:kube-system::extension-apiserver-authentication
```

## Practice2: Working with pod manifest files

Note: Try not to do a copy/paste on commands requests unless you are instructed to do so. Copy/paste will not help you to learn Kubernetes!

1. Now it is time to deploy pod using manifest file (declarative approach). Copy the following code block on your local computer in a file called redis.yaml:

```
apiVersion: v11
kind: pod
metadata:
  name: static-web
labels:
  role: myrole
specs:
containers:
- name: redis
  image: redis123
```

2. Try to deploy the pod defined in redis.yaml. Run **kubectl create -f redis.yaml**.

```
PS C:\Users\V&M\Desktop\lab> kubectl create -f redis.yaml
error: resource mapping not found for name: "" namespace: "" from "redis.yaml": no matches for kind "pod" in version "v11"
ensure CRDs are installed first
PS C:\Users\V&M\Desktop\lab>
```

3. You will receive errors on your screen. Your next task will be to correct the syntax of the code you just copied. You can use the online Kubernetes documentation or you can search the internet in general.

- v11 is not a valid API version. The correct API version for creating Pods is v1
- specs should be spec
- wrong indentation

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: static-web
5  labels:
6    role: myrole
7  spec:
8    containers:
9    - name: redis
10     image: redis123
```

4. When you solve all the syntax errors your pod should be deployed but is it running? What is the status of your pod? **The status is ImagePullBackOff.**

```
PS C:\Users\V&M\Desktop\lab> kubectl create -f redis.yaml
pod/static-web created
PS C:\Users\V&M\Desktop\lab> kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE
static-web    0/1     ImagePullBackOff   0          27s
PS C:\Users\V&M\Desktop\lab>
```

5. Check the events associated with this pod. Run the **kubectl describe pod static-web** command. What are the events showing? Why your pod is not running?

**The error messages indicates that the image redis123 could not be pulled from the specified container registry.**

```
Events:
  Type     Reason      Age          From          Message
  ----     -
  Normal   Scheduled   93s         default-scheduler   Successfully assigned default/static-web to docker-desktop
  Warning   Failed      45s (x3 over 91s) kubelet        Failed to pull image "redis123": rpc error: code = Unknown desc = Error response from daemon: pull access denied for redis123, repository does not exist or may require 'docker login': denied: requested access to the resource is denied
  Warning   Failed      45s (x3 over 91s) kubelet        Error: ErrImagePull
  Normal   BackOff     15s (x4 over 90s) kubelet        Back-off pulling image "redis123"
  Warning   Failed      15s (x4 over 90s) kubelet        Error: ImagePullBackOff
  Normal   Pulling     2s (x4 over 93s) kubelet        Pulling image "redis123"
```

6. Find the correct image (check the Docker hub page) and correct it in the manifest.

```
apiVersion: v1
kind: Pod
metadata:
  name: static-web
  labels:
    role: myrole
spec:
  containers:
  - name: redis
    image: asterixlegaulois/redis123
```

7. Locate the image information and put the correct image name. Redeploy the pod (first run **kubectl delete pod static-web** to delete the pod, then run **kubectl create** once again).

```
PS C:\Users\V&M\Desktop\lab> kubectl delete pod static-web
pod "static-web" deleted
PS C:\Users\V&M\Desktop\lab> kubectl create -f redis.yaml
pod/static-web created
PS C:\Users\V&M\Desktop\lab> kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
static-web    1/1     Running   0          7s
PS C:\Users\V&M\Desktop\lab>
```

8. Check the status of your pod. It should be running now. **Running**

9. Now you can delete the pod. Try to delete it using the **kubectl delete -f redis.yaml**.

```
PS C:\Users\V&M\Desktop\lab> kubectl delete -f redis.yaml
pod "static-web" deleted
PS C:\Users\V&M\Desktop\lab>
```

10. Your next task is to create and test nginx pod definition. Your definition should use the nginx official image, should use label named app with value frontend and should publish port 80. Make sure you complete this task because we will use this template in our next Labs. Your nginx pod should be running without any issues.

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: frontend
spec:
  containers:
  - name: nginx-container
    image: nginx
    ports:
    - containerPort: 80
```

```
Events:
  Type     Reason      Age   From              Message
  ----     -
  Normal   Scheduled   77s   default-scheduler Successfully assigned default/nginx-pod to docker-desktop
  Normal   Pulling     77s   kubelet           Pulling image "nginx"
  Normal   Pulled      63s   kubelet           Successfully pulled image "nginx" in 13.57627594s
  Normal   Created     63s   kubelet           Created container nginx-container
  Normal   Started     63s   kubelet           Started container nginx-container
```

11. Final task of this practice will be to define pod definition with following details:

- Image=memcached
- Port= 11211
- Label app=web
- CPU request=0.35 cores
- RAM request=0.15 GB
- CPU limit=0.5 cores
- Ram limit=0.25 GB
- Restart policy=Never

```
apiVersion: v1
kind: Pod
metadata:
  name: memcached-web
  labels:
    app: web
spec:
  containers:
  - name: memcached
    image: memcached
    ports:
    - containerPort: 11211
    resources:
      requests:
        cpu: "0.35"
        memory: "0.15Gi"
      limits:
        cpu: "0.5"
        memory: "0.25Gi"
    imagePullPolicy: IfNotPresent
    restartPolicy: Never
```



12. Don't forget to try your pod definition.

```
PS C:\Users\V&M\Desktop\lab> kubectl create -f memcached.yaml
pod/memcached-web created
PS C:\Users\V&M\Desktop\lab> kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
memcached-web 1/1     Running   0           15s
nginx-pod     1/1     Running   0           7m7s
PS C:\Users\V&M\Desktop\lab>
```

## Practice3: Multi-container pods

Note: Try not to do a copy/paste on commands requests unless you are instructed to do so.

Copy/paste will not help you to learn Kubernetes!

1. Once finished you can try to create multi-container pod definition. Your multi-container pod should use redis and nginx containers with port 6379 and 80 published respectively. Label name should be app with value web.

```
apiVersion: v1
kind: Pod
metadata:
  name: webapp
  labels:
    app: web
spec:
  containers:
    - name: redis
      image: asterixlegaulois/redis123
      ports:
        - containerPort: 6379
    - name: nginx-container
      image: nginx
      ports:
        - containerPort: 80
```

2. Note that in reality there is no sense to put the redis and nginx under the same pod but it can be done for the purpose of learning.

3. Deploy your multi-container pod. It should have running status. What is written under Ready column when you kubectl get the pods? Why your pod displays different values for ready? Under Ready column is written 2/2 because in this pod we have two containers.

```
PS C:\Users\V&M\Desktop\lab> kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
memcached-web 1/1     Running   0           16m
nginx-pod     1/1     Running   0           28m
webapp        2/2     Running   0           66s
PS C:\Users\V&M\Desktop\lab>
```



4. Kubectl describe your new pod, and locate the containers section. How many containers are listed? **Two**.

```
Containers:
  redis:
    Container ID:   docker://e2630b2c8b28c2339972f5dfb8ddf6d77e9cbbcc36e908e59bc1e4238589e58a
    Image:          asterixlegaulois/redis123
    Image ID:       docker-pullable://asterixlegaulois/redis123@sha256:cb2ddf11373c66e8c66ea7cb4
ddd82bd69849d4838ae41ee0e71a0bc5c6cc4c4
    Port:          6379/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Tue, 04 Apr 2023 13:08:44 +0200
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-8zgfw (ro)
  nginx-container:
    Container ID:   docker://12a72152b265220b56f0bc81f61f78265c0c97a76a53c04a94aea93ee4833f5a
    Image:          nginx
    Image ID:       docker-pullable://nginx@sha256:2ab30d6ac53580a6db8b657abf0f68d75360ff5cc1670
a85acb5bd85ba1b19c0
    Port:          80/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Tue, 04 Apr 2023 13:08:45 +0200
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-8zgfw (ro)
```

5. Delete all the pods under the default namespace.

```
PS C:\Users\V&M\Desktop\lab> kubectl delete pods --all
>>
pod "memcached-web" deleted
pod "nginx-pod" deleted
pod "webapp" deleted
PS C:\Users\V&M\Desktop\lab> _
```

6. Don't delete any of the manifest files you have created so far.

## Practice4: Probes

Note: Try not to do a copy/paste on commands requests unless you are instructed to do so.

Copy/paste will not help you to learn Kubernetes!

1. First we will create and test liveness probe with exec test. Create a file named probes\_exec.yaml with following content:

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
name: liveness-exec
spec:
  containers:
  - name: liveness
    image: k8s.gcr.io/busybox
    args:
    - /bin/sh
    - -c
    - touch /tmp/healthy; sleep 30; rm -rf /tmp/healthy; sleep 600
  livenessProbe:
    exec:
      command:
      - cat
      - /tmp/healthy
    initialDelaySeconds: 5
    periodSeconds: 5
```

2. Examine the containers args commands especially the line that start with touch. This bash pipeline will help us to test the liveness probes.

The container runs a shell command that touches a file named /tmp/healthy, sleeps for 30 seconds, removes the file and then sleeps for 600 seconds.

3. Run **kubectl create -f probes\_exec.yaml**.

```
PS C:\Users\V&M\Desktop\lab> kubectl create -f probes_exec.yaml
pod/liveness-exec created
```

4. Run **kubectl describe pod liveness-exec** immediately after you deploy the pod. The output should indicate that no liveness probes have failed yet.

```
Events:
  Type     Reason      Age   From              Message
  ----     -
  Normal   Scheduled   15s   default-scheduler Successfully assigned default/liveness-exec to doc
ker-desktop
  Normal   Pulling     14s   kubelet           Pulling image "k8s.gcr.io/busybox"
  Normal   Pulled      12s   kubelet           Successfully pulled image "k8s.gcr.io/busybox" in
2.11243698s
  Normal   Created     12s   kubelet           Created container liveness
  Normal   Started     12s   kubelet           Started container liveness
```

- After 35 seconds, view the Pod events again. Run **kubectl describe pod liveness-exec**.
- At the bottom of the output, there should be a messages indicating that the liveness probes have failed, and the containers have been killed and recreated.

```
Events:
  Type       Reason            Age           From          Message
  ----       -
  Normal     Scheduled         68s          default-scheduler   Successfully assigned default/liveness-exec to docker-desktop
  Normal     Pulling           67s          kubelet         Pulling image "k8s.gcr.io/busybox"
  Normal     Pulled            65s          kubelet         Successfully pulled image "k8s.gcr.io/busybox" in 2.11243698s
  Normal     Created           65s          kubelet         Created container liveness
  Normal     Started           65s          kubelet         Started container liveness
  Warning    Unhealthy         23s (x3 over 33s) kubelet         Liveness probe failed: cat: can't open '/tmp/healthy': No such file or directory
  Normal     Killing           23s          kubelet         Container liveness failed liveness probe, will be restarted
```

- Wait another 30 seconds, and verify that the container has been restarted. Run **kubectl get pod liveness-exec**.
- The output should show that RESTARTS has been incremented.

```
PS C:\Users\V&M\Desktop\lab> kubectl get pod liveness-exec
NAME          READY   STATUS    RESTARTS   AGE
liveness-exec 1/1     Running   1 (55s ago) 2m10s
PS C:\Users\V&M\Desktop\lab> _
```

- We will continue with HTTP probe. Create file named probes\_http.yaml with following content:

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
  name: liveness-http
spec:
  containers:
    - name: liveness
      image: k8s.gcr.io/liveness
      args:
        - /server
      livenessProbe:
        httpGet:
          path: /healthz
          port: 8080
          httpHeaders:
            - name: Custom-Header
              value: Awesome
        initialDelaySeconds: 3
        periodSeconds: 3
```

10. Just for your info, /healthz handler has following function implemented:

```
http.HandleFunc("/healthz", func(w http.ResponseWriter, r *http.Request) {
    duration := time.Now().Sub(started)
    if duration.Seconds() > 10 {
        w.WriteHeader(500)
        w.Write([]byte(fmt.Sprintf("error: %v", duration.Seconds())))
    } else {
        w.WriteHeader(200)
        w.Write([]byte("ok"))
    }
})
```

11. For the first 10 seconds that the container is alive, the /healthz handler returns a status of 200. After that, the handler returns a status of 500.

12. Run **kubectl create -f probes\_http.yaml**.

13. Immediately run (you only have 10 secs to run this command) **kubectl describe pod liveness-http**.

14. Your pod should be live and running.

```
PS C:\Users\V&M\Desktop\lab> kubectl create -f probes_http.yaml
pod/liveness-http created
PS C:\Users\V&M\Desktop\lab> kubectl describe pod liveness-http
Name:          liveness-http
Namespace:     default
Priority:       0
Service Account: default
Node:          docker-desktop/192.168.65.4
Start Time:    Tue, 04 Apr 2023 13:37:15 +0200
Labels:        test=liveness
Annotations:    <none>
Status:        Running
IP:            10.1.0.18
```

15. After 10 seconds, view Pod events to verify that liveness probes have failed and the container has been restarted. Run again **kubectl describe pod liveness-http**.

```
Events:
  Type     Reason      Age           From              Message
  ----     -
  Normal   Scheduled   80s           default-scheduler  Successfully assigned default/liveness-http to docker-desktop
  Normal   Pulled      78s           kubelet           Successfully pulled image "k8s.gcr.io/liveness" in 1.216349349s
  Normal   Pulled      60s           kubelet           Successfully pulled image "k8s.gcr.io/liveness" in 1.144994584s
  Normal   Created    42s (x3 over 78s) kubelet           Created container liveness
  Normal   Started    42s (x3 over 78s) kubelet           Started container liveness
  Normal   Pulled     42s           kubelet           Successfully pulled image "k8s.gcr.io/liveness" in 1.182641358s
  Normal   Pulling    26s (x4 over 79s) kubelet           Pulling image "k8s.gcr.io/liveness"
  Warning  Unhealthy  26s (x9 over 68s) kubelet           Liveness probe failed: HTTP probe failed with statuscode: 500
  Normal   Killing    26s (x3 over 62s) kubelet           Container liveness failed liveness probe, will be restarted
```

16. You should see the same output as in step 7. Kubelet will reboot the container.

17. We continue with TCP probes. Create file named probes\_tcp.yaml with following content:

```
apiVersion: v1
kind: Pod
metadata:
  name: liveness-tcp
  labels:
    app: goproxy
spec:
  containers:
  - name: goproxy
    image: k8s.gcr.io/goproxy:0.1
    ports:
    - containerPort: 8080
    livenessProbe:
      tcpSocket:
        port: 9999 #8080 is valid port
      initialDelaySeconds: 15
      periodSeconds: 20
```

18. Run **kubectl create -f probes\_tcp.yaml**.

19. Immediately run (you only have 10 secs to run this command) **kubectl describe pod liveness-tcp**.

20. Your pod should be live and running.

```
PS C:\Users\V&M\Desktop\lab> kubectl create -f probes_tcp.yaml
pod/liveness-tcp created
PS C:\Users\V&M\Desktop\lab> kubectl describe pod liveness-tcp
Name:          liveness-tcp
Namespace:     default
Priority:       0
Service Account: default
Node:          docker-desktop/192.168.65.4
Start Time:    Tue, 04 Apr 2023 13:58:20 +0200
Labels:        app=goproxy
Annotations:    <none>
Status:        Running
IP:            10.1.0.19
```

21. After 10 seconds, view Pod events to verify that liveness probes have failed and the container has been restarted. Run again **kubectl describe pod liveness-tcp**.

```
Events:
  Type     Reason      Age           From              Message
  ----     -
  Normal    Scheduled   74s          default-scheduler  Successfully assigned default/liveness-tcp to docker-desktop
  Normal    Pulling     73s          kubelet           Pulling image "k8s.gcr.io/goproxy:0.1"
  Normal    Pulled      71s          kubelet           Successfully pulled image "k8s.gcr.io/goproxy:0.1" in 2.233846916s
  Normal    Created     13s (x2 over 70s) kubelet           Created container goproxy
  Normal    Started     13s (x2 over 70s) kubelet           Started container goproxy
  Warning   Unhealthy   13s (x3 over 53s) kubelet           Liveness probe failed: dial tcp 10.1.0.19:9999: connect: connection refused
  Normal    Killing     13s          kubelet           Container goproxy failed liveness probe, will be restarted
```

22. You should see the same output as in step 7 and 16. Kubelet will reboot the container.

23. Our last job will be to define one readiness probe using HTTP test.  
24. Create file named readiness\_http.yaml with following content:

```
apiVersion: v1
kind: Pod
metadata:
  name: readiness-http
  labels:
  app: test
spec:
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 80
    readinessProbe:
      initialDelaySeconds: 1
      periodSeconds: 2
      timeoutSeconds: 1
      successThreshold: 1
      failureThreshold: 1
      httpGet:
        host:
        scheme: HTTP
        path: /
        httpHeaders:
        - name: Host
          value: myapplication1.com
        port: 80
```

25. Run `kubect1 create -f readiness_http.yaml`.  
26. Run `kubect1 get pods -A` to see the status of your pod.  
27. Pods and their status and ready states will be displayed; our pod should be in running state.

```
PS C:\Users\V&M\Desktop\lab> kubect1 create -f readiness_http.yaml
pod/readiness-http created
PS C:\Users\V&M\Desktop\lab> kubect1 get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	liveness-exec	0/1	CrashLoopBackOff	13 (56s ago)	37m
default	liveness-http	0/1	CrashLoopBackOff	13 (3m46s ago)	28m
default	liveness-tcp	0/1	CrashLoopBackOff	5 (70s ago)	7m11s
default	readiness-http	1/1	Running	0	7s
kube-system	coredns-565d847f94-f45t7	1/1	Running	0	120m
kube-system	coredns-565d847f94-sjcxj	1/1	Running	0	120m
kube-system	etcd-docker-desktop	1/1	Running	1	120m
kube-system	kube-apiserver-docker-desktop	1/1	Running	1	120m
kube-system	kube-controller-manager-docker-desktop	1/1	Running	1	120m
kube-system	kube-proxy-6h7fw	1/1	Running	0	120m
kube-system	kube-scheduler-docker-desktop	1/1	Running	1	120m
kube-system	storage-provisioner	1/1	Running	0	119m
kube-system	vpnkit-controller	1/1	Running	9 (7m8s ago)	119m

28. Run **kubectl describe pod readiness-http**. Examine the events for this pod. Everything should be OK.

```
Events:
  Type       Reason      Age   From                Message
  ----       -
  Normal     Scheduled   110s  default-scheduler   Successfully assigned default
/pod/readiness-http to docker-desktop
  Normal     Pulling     109s  kubelet             Pulling image "nginx"
  Normal     Pulled      107s  kubelet             Successfully pulled image "ng
inx" in 1.48232177s
  Normal     Created     107s  kubelet             Created container nginx
  Normal     Started     107s  kubelet             Started container nginx
```

29. Now delete the pod and edit the readiness\_http.yaml so that the port parameter has 81 value.

```
PS C:\Users\V&M\Desktop\lab> kubectl delete pod readiness-http
pod "readiness-http" deleted
```

```
readinessProbe:
  initialDelaySeconds: 1
  periodSeconds: 2
  timeoutSeconds: 1
  successThreshold: 1
  failureThreshold: 1
  httpGet:
    host:
    scheme: HTTP
    path: /
    httpHeaders:
    - name: Host
      value: myapplication1.com
    port: 81
```

30. Run again **kubectl create -f readiness\_http.yaml**.

31. Run **kubectl get pods -A** to see the status of your pod. You should see that the pod is running but it is not in ready state.

```
PS C:\Users\V&M\Desktop\lab> kubectl create -f readiness_http.yaml
pod/readiness-http created
PS C:\Users\V&M\Desktop\lab> kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	liveness-exec	0/1	CrashLoopBackOff	13 (4m46s ago)	41m
default	liveness-http	0/1	CrashLoopBackOff	15 (111s ago)	32m
default	liveness-tcp	0/1	CrashLoopBackOff	7 (80s ago)	11m
default	readiness-http	0/1	Running	0	10s
kube-system	coredns-565d847f94-f45t7	1/1	Running	0	123m
kube-system	coredns-565d847f94-sjcxj	1/1	Running	0	123m
kube-system	etcd-docker-desktop	1/1	Running	1	123m
kube-system	kube-apiserver-docker-desktop	1/1	Running	1	124m
kube-system	kube-controller-manager-docker-desktop	1/1	Running	1	123m
kube-system	kube-proxy-6h7fw	1/1	Running	0	123m
kube-system	kube-scheduler-docker-desktop	1/1	Running	1	124m
kube-system	storage-provisioner	1/1	Running	0	123m
kube-system	vpnkit-controller	1/1	Running	10 (2m55s ago)	123m



32. Describe the pod. Run **kubectl describe pod readiness-http**.

33. From the events we can see that readiness probe failed due to the connection being refused therefore pod will not receive any traffic.

```
Events:
  Type      Reason      Age           From          Message
  ----      -
  Normal    Scheduled   55s           default-scheduler Successfully assigned de
fault/readiness-http to docker-desktop
  Normal    Pulling     54s           kubelet        Pulling image "nginx"
  Normal    Pulled      53s           kubelet        Successfully pulled imag
e "nginx" in 1.48674243s
  Normal    Created     53s           kubelet        Created container nginx
  Normal    Started     52s           kubelet        Started container nginx
  Warning   Unhealthy   13s (x21 over 51s) kubelet        Readiness probe failed:
Get "http://10.1.0.21:81/": dial tcp 10.1.0.21:81: connect: connection refused
```

34. Delete all pods under the default namespace.

```
PS C:\Users\V&M\Desktop\lab> kubectl delete pods --all -n default
>>
pod "liveness-exec" deleted
pod "liveness-http" deleted
pod "liveness-tcp" deleted
pod "readiness-http" deleted
PS C:\Users\V&M\Desktop\lab>
```

35. Don't delete any manifest files created so far.